

Scaling Biclique Percolation for Community Detection in Large Bipartite Networks

Yue Zeng

Beijing Institute of Technology
bruceez@163.com

Rong-Hua Li

Beijing Institute of Technology
lironghuabit@126.com

Qiangqiang Dai

Beijing Institute of Technology
qiangd66@gmail.com

Guoren Wang

Beijing Institute of Technology
wanggrbit@126.com

(4,4)-BCPC




Fig. 1. A (4,4)-BCPC around Jiawei Han in the DBLP author-paper bipartite graph

Abstract—Community detection is a core network-analysis task, and many real-world networks are naturally bipartite. Biclique percolation community (BCPC) detects cohesive bipartite structures by connecting maximal bicliques that share an (α, β) -biclique. Existing BCPC methods are costly because they explicitly traverse the potentially massive maximal biclique adjacency graph (MBAG). We first propose a cross-subtree partial-BCPC based solution: maximal bicliques that provably belong to the same BCPC, even when scattered across different enumeration subtrees, are grouped as one MBAG vertex, greatly reducing graph size. We further go beyond MBAG traversal with an (α, β) -biclique enumeration approach, which connects maximal bicliques sharing an enumerated (α, β) -biclique and again uses cross-subtree partial-BCPCs to prune the enumeration space. Experiments show that our methods outperform existing methods by nearly three orders of magnitude.

I. INTRODUCTION

Community detection aims to uncover closely connected vertex groups and is central to network analysis [1], with applications in social networks [2], biological systems [3], and recommendation systems [4]. Many real-world networks are bipartite graphs $G = (U, V, E)$, where vertices are split into two disjoint sets and edges only connect the two sides, e.g., user-item networks [5] and author-paper networks [6]. This structure motivates specialized community models for bipartite graphs.

The biclique percolation community (BCPC) is a widely studied approach for detecting cohesive structures in bipartite graphs [7], [8], [9], [10], [11], [12], [13], [14]. A biclique is a complete bipartite subgraph, and a biclique percolation community is defined as a maximal set of bicliques that are connected through *adjacency*. In the (α, β) -biclique percolation community, two bicliques are considered *adjacent* if they share at least α vertices on one side and β on the other (share an (α, β) -biclique).

Example 1. Figure 1 shows a (4,4)-BCPC around Jiawei Han in the DBLP author-paper bipartite graph. The dashed

boxes B_1, B_2, B_3 are three maximal bicliques, none of which individually captures the full community. Since B_1-B_2 and B_2-B_3 each share a (4,4)-biclique, they are linked by biclique adjacency and percolate into a single BCPC.

As illustrated in Example 1, a BCPC is formed by multiple maximal bicliques that *percolate* into a single community without requiring the whole to be a complete biclique, making it a natural relaxation of single-biclique models. The parameters α, β act as cohesion controls: smaller values yield larger, looser communities, while larger values yield tighter, higher-confidence ones.

BCPC has been applied to social networks for analyzing user roles and themes [11], Wikipedia editor-article networks for revealing collaborative editing clusters [8], online-course resource networks for mining learners' access and participation patterns [10], [12], and enterprise security graphs for supporting host-user-application monitoring [9]. BCPC naturally allows entities to participate in multiple communities, which better matches real-world bipartite networks where entities often play multiple roles or appear in multiple functional contexts [11], [8], [10], [9].

Despite its widespread practical utility, the BCPC method faces significant computational challenges due to its high time complexity. To the best of our knowledge, the existing BCPC methods are all based on [7]. Its framework was later implemented in [15] and used to build indexes for personalized search problems [13], [14]. The basic idea is based on maximal biclique adjacency graph (MBAG). This method first enumerates all maximal bicliques and then uses them as vertices to construct a maximal biclique adjacency graph. In this graph, there is an edge between two vertices if and only if the overlapping part of the corresponding two maximal bicliques contains an (α, β) -biclique. The connected components in the adjacency graph correspond one-to-one with the (α, β) -BCPC. However, real-world bipartite graphs may contain a large number of maximal bicliques (up to $2^{n/2}$ theoretically [16]), and the adjacency graph formed by these maximal bicliques could contain a significant number of edges (up to 10^{11} , see Table II). Traversing such a large-scale adjacency graph is highly time-consuming.

To overcome these limitations, we first reduce the MBAG through partial-BCPCs. In the maximal biclique enumeration tree, if a prefix already contains at least α vertices from U and β vertices from V , all maximal bicliques under this branch share an (α, β) -biclique and belong to the same BCPC; they can therefore be represented as one MBAG vertex. We further develop a *cross-subtree* compression strategy: using local

information at search tree nodes, it safely merges maximal bicliques that share the same (α, β) -biclique but are scattered across different subtrees. Stop-labels further let the search terminate once a subtree can be represented by an existing partial-BCPC label. This retains as few as 2.7% of the original MBAG vertices on Youtube with $\alpha = \beta = 2$ (Exp-3).

Since the partial-BCPC solution still traverses a reduced MBAG, we also propose a second approach based on (α, β) -biclique enumeration. It enumerates (α, β) -bicliques and uses union-find to connect all maximal bicliques containing each (α, β) -biclique, exactly matching the adjacency condition. Cross-subtree partial-BCPCs significantly prune the exponential search [17], [18]: a branch is discarded if its intermediate biclique is contained by no maximal biclique, or only by maximal bicliques already in one partial-BCPC. This reduces visited enumeration nodes by over three orders of magnitude (Exp-4) and yields over $390\times$ speedup over the SOTA MBAG method (Exp-1, Youtube, $\alpha = \beta = 2$).

We summarize our contributions as follows.

Novel partial-BCPC based solution. We introduce partial-BCPCs, a relaxed BCPC notion, to reduce the MBAG. Besides within-subtree merging on the maximal biclique enumeration tree, we develop a provably correct *cross-subtree* strategy that groups maximal bicliques sharing the same (α, β) -biclique even when they are scattered across different subtrees, with stop-labels reducing the search overhead. This retains as few as 2.7% of MBAG vertices (Exp-3, Youtube, $\alpha = \beta = 2$).

Novel (α, β) -biclique enumeration based solution. We also propose a framework that directly enumerates (α, β) -bicliques and connects the maximal bicliques sharing them. Partial-BCPC pruning reduces the enumeration tree by up to three orders of magnitude (Exp-4).

Extensive experiments. On 10 real-world datasets, our algorithms achieve large speedups. For example, on Youtube with $\alpha = \beta = 2$, MBAG takes 13,376 seconds, whereas our best algorithm BicliqueP takes only 34 seconds. A case study shows that BCPC captures more meaningful communities than biclique, bitruss, and bicore.

Reproducibility. The source code of this paper is released on <https://github.com/bruce1114/bcpc> for reproducibility purposes.

II. PRELIMINARIES

Let $G = (U, V, E)$ be an undirected and unweighted bipartite graph, where U, V are two disjoint sets of vertices and E is a set of edges in form of $(u, v), u \in U, v \in V$. For a vertex $u \in U$, $Nei_G(u)$ is the set of neighbors of u , i.e., $Nei_G(u) = \{v | (u, v) \in E\}$. A bipartite graph $G' = (U', V', E')$ is a subgraph of G if $U' \subseteq U, V' \subseteq V, E' \subseteq E$.

Definition 1 (Biclique). Given a bipartite graph $G = (U, V, E)$, a biclique is a pair of vertex sets $B = (X, Y)$ where $X \subseteq U, Y \subseteq V$ and $\forall u \in X, v \in Y, (u, v) \in E$.

For simplicity, we let $B.X = X, B.Y = Y$ if $B = (X, Y)$. For two bicliques $B, B', B \subseteq B'$ indicates that $B.X \subseteq B'.X, B.Y \subseteq B'.Y$. A maximal biclique B is a biclique that there is no other biclique B' satisfying $B \subseteq B'$.




Fig. 2. A bipartite graph G and (α, β) -BCPC ($\alpha = 2, \beta = 2$)

TABLE I
FREQUENTLY USED NOTATIONS

Notation	Description
$G = (U, V, E)$	bipartite graph with vertex sets U, V and edge set E
$B = (B.X, B.Y)$	a biclique with $B.X \subseteq U$ and $B.Y \subseteq V$
$\mathcal{B}$	an (α, β) -BCPC, i.e., a set of maximal bicliques
$\mathcal{P}$	a partial-BCPC, i.e., a subset of some $\mathcal{B}$
MBAG	maximal biclique adjacency graph
MBE tree	maximal biclique enumeration tree
nd	a node in the MBE tree
$nd.R_U, nd.R_V$	current biclique vertices stored at nd
$nd.C_U, nd.C_V$	candidate vertices for expansion at nd
$nd.X_U, nd.X_V$	exclusion vertices at nd
$nd.RC X_U$	shorthand for $nd.R_U \cup nd.C_U \cup nd.X_U$; similarly for $nd.RC X_V$
(α, β) -node	node where $ R_U \geq \alpha$ and $ R_V \geq \beta$ first hold

Definition 2 ((α, β) -Biclique Adjacency). Given two integers α, β and two bicliques B_1, B_2 , B_1 and B_2 are (α, β) -adjacent if $|B_1.X \cap B_2.X| \geq \alpha, |B_1.Y \cap B_2.Y| \geq \beta$.

Definition 3 ((α, β) -Biclique Connectivity). Given two integers α, β and two bicliques B_1, B_n , B_1 and B_n are (α, β) -connected if there is a sequence of bicliques $B_1, B_2, \dots, B_n$ that $\forall i = 1, \dots, n-1, B_i$ and B_{i+1} are (α, β) -adjacent.

Definition 4 ((α, β) -Biclique Percolation Community ((α, β) -BCPC)). Given a bipartite graph G and two integers α, β , an (α, β) -biclique percolation community is a set $\mathcal{B}$ of maximal bicliques in G satisfying the following two conditions: (i) any two maximal bicliques in $\mathcal{B}$ are (α, β) -connected; and (ii) $\mathcal{B}$ is maximal, i.e., no maximal biclique outside $\mathcal{B}$ is (α, β) -connected with any maximal biclique in $\mathcal{B}$ (equivalently, $\mathcal{B}$ can not be enlarged by adding any other maximal biclique while keeping condition (i)).

Example 2. Let $\alpha = 2, \beta = 2$, Figure 2 illustrates a bipartite graph G and some of the $(2, 2)$ -BCPCs. These BCPCs locate on the left, in the middle, and on the right, respectively. The $(2, 2)$ -BCPC in the middle is composed of three maximal bicliques: $B_1 = (\{u_1, u_2\}, \{v_1, v_2, v_3\})$, $B_2 = (\{u_1, u_2, u_3\}, \{v_2, v_3\})$, and $B_3 = (\{u_2, u_3\}, \{v_2, v_3, v_4\})$. Among these bicliques, B_1 and B_2 are adjacent because their intersection, $(\{u_1, u_2\}, \{v_2, v_3\})$, satisfies the condition in Definition 2. For the same reason, B_2 and B_3 are also adjacent. Therefore, B_1, B_2 , and B_3 can together form a valid $(2, 2)$ -BCPC. On the left and right sides of Figure 2 are two maximal bicliques. These maximal bicliques are not adjacent (Definition 2) to any other maximal bicliques, and thus can be regarded as two special BCPCs.

Problem statement. Given a bipartite graph G and two integers α, β , the goal of our work is to detect all (α, β) -BCPCs (BCPC detection).

Algorithm 1: The basic maximal biclique enumeration framework in [16]

Input: Bipartite graph $G = (U, V, E)$
Output: All maximal bicliques in G

```

1 Enum  $(\emptyset, \emptyset, U, V, \emptyset, \emptyset)$ ;
2 Procedure Enum  $(R_U, R_V, C_U, C_V, X_U, X_V)$ 
3 if  $C_U \cup C_V = \emptyset$  then
4   if  $X_U \cup X_V = \emptyset$  then Output  $(R_U, R_V)$  as a maximal biclique;
5   return;
6 foreach  $u \in C_U$  do
7   Branch  $(R_U, R_V, u, C_U - \{u\}, C_V, X_U, X_V)$ ;
8    $C_U \leftarrow C_U - \{u\}$ ;  $X_U \leftarrow X_U \cup \{u\}$ ;
9 foreach  $v \in C_V$  do
10  Branch  $(R_V, R_U, v, C_V - \{v\}, C_U, X_V, X_U)$ ;
11   $C_V \leftarrow C_V - \{v\}$ ;  $X_V \leftarrow X_V \cup \{v\}$ ;
12 Procedure Branch  $(R_U, R_V, u, C_U, C_V, X_U, X_V)$ 
13   $C'_U \leftarrow C_U \cap \text{Nei}_G(u)$ ;  $X'_U \leftarrow X_U \cap \text{Nei}_G(u)$ ;
14  Enum  $(R_U \cup \{u\}, R_V, C'_U, C_V, X'_U, X_V)$ ;

```

A. Maximal Biclique Enumeration

The existing state-of-the-art maximal biclique enumeration methods [19], [20], [16] are all based on the set-enumeration tree structure, where vertices are iteratively traversed to build nodes of the enumeration tree, with each node representing a distinct biclique. In this paper, we adopt the algorithmic framework in [16] to perform maximal biclique enumeration and to design the partial-BCPC based method in Section IV. This is because the method in [16] achieves state-of-the-art performance in maximal biclique enumeration, and furthermore, the framework in [16] is adaptable to a wider range of cohesive subgraph models, making our partial-BCPC based method in Section IV extensible to other cohesive subgraph models in bipartite graphs.

The basic framework [16]. Algorithm 1 presents the basic framework for maximal biclique enumeration in [16]. This framework recursively explores candidate vertex sets to construct the set-enumeration tree, while leveraging exclusion sets to avoid non-maximal bicliques. Algorithm 1 begins by invoking the recursive procedure Enum in Line 1. The parameters of Enum contain the current biclique (R_U, R_V) , candidate vertices (C_U, C_V) for expansion, and exclusion vertices (X_U, X_V) that have already been processed. Within Enum, the algorithm first checks the termination condition: $C_U \cup C_V = \emptyset$ (Line 3). If the termination condition is satisfied while $X_U \cup X_V = \emptyset$ (Line 4), it implies that (R_U, R_V) is a maximal biclique. Subsequently, the algorithm iterates over each vertex u in the candidate set C_U (Line 6), calling the Branch procedure to perform branch expansion (Line 7), and moves u from the candidate set C_U to the exclusion set X_U (Line 8). Symmetrically, the algorithm performs the same operation for each vertex v in C_V (Lines 9-11). In the Branch procedure, the algorithm computes the new candidate set $C'_U = C_U \cap \text{Nei}_G(u)$ and the new exclusion set $X'_U = X_U \cap \text{Nei}_G(u)$, adds u to R_U , and recursively calls the Enum procedure to further explore maximal bicliques within this branch.

Pivoting technique [16]. A notable issue with the aforementioned framework is that it generates a significant amount of unnecessary computations that produce non-maximal bicliques. We first make the following observation: for any vertex $u \in U$ in a bipartite graph $G = (U, V)$, and for any maximal biclique $B = (X, Y)$ in the bipartite graph, exactly




Fig. 3. The partial enumeration tree of the framework in [16] on bipartite graph of Figure 2. Pivot vertices are enclosed in dashed circles. $nd(B_i, h_j)$ represents the tree node at level h_j on the B_i branch of the enumeration tree one of the following two conditions must hold:

- 1) B contains u (i.e., $u \in X$).
- 2) B contains a vertex v , where v is a non-neighbor of u in V (i.e., $\exists v \in V - \text{Nei}_G(u), v \in Y$).

Based on the above observation, we can conclude that by selecting a pivot vertex u from X_U , C_U , X_V , or C_V , and enumerating only u (if $u \notin X_U, X_V$) and vertices in C_V (or C_U , if $u \in X_V, C_V$) that are not neighbors of u in Lines 6-11 of Algorithm 1, we can ensure the completeness of all maximal bicliques while reducing unnecessary enumeration branches that produce non-maximal bicliques. An effective method for selecting a pivot vertex is to choose the vertex in $X_U \cup C_U$ (or $X_V \cup C_V$) that has the fewest non-neighbor vertices in C_V (or C_U) as the pivot vertex.

Vertex ordering [20]. The efficiency of maximal biclique enumeration algorithms is often influenced by the order in which vertices are processed. [16] applied the 2-hop degree ordering introduced in [20] to their framework. This vertex ordering is only applied at the first level of the recursive search, while the pivoting pruning technique mentioned earlier is utilized in subsequent levels of the recursive search.

Example 3. Figure 3 illustrates a partial enumeration tree of the maximal biclique enumeration framework [16] on the bipartite graph in Figure 2. The root node initializes its candidate set C_U with the vertices in U ordered by 2-hop degree ordering, while only the branches starting from u_3 and u_2 are shown. For example, after selecting u_3 , the node $nd(B_1, h_1)$ adds u_3 to R_U , updates C_V to $\text{Nei}_G(u_3)$, and records previously processed vertices u_4, u_5, u_6 in X_U .

At $nd(B_1, h_1)$, the framework selects v_3 as the pivot and only expands the branch corresponding to v_3 , skipping vertices that cannot lead to new maximal bicliques under the pivoting rule. After v_3 is added to R_V , the candidate and exclusion sets on the U -side are updated by intersecting their old values with $\text{Nei}_G(v_3)$. When both candidate sets C_U and C_V become empty, and the exclusion sets X_U and X_V are also empty, the current pair (R_U, R_V) is output as a maximal biclique.

B. (α, β) -Biclique Enumeration

The existing state-of-the-art (α, β) -biclique enumeration method is proposed in [17]. Given a bipartite graph $G =$

Algorithm 2: The (α, β) -biclique enumeration framework in [17]

Input: Bipartite graph $G = (U, V, E)$; Two integers α, β
Output: All (α, β) -bicliques in G

- 1 Construct 2-hop graph $\vec{H} = (U, \vec{E}_h)$ on vertex set U [17];
/ $\vec{H} = (U, \vec{E}_h)$: a directed graph, $(u, v) \in \vec{E}_h$ indicates that $u < v$, u, v have common neighbors in V */*
- 2 BicliqueListing($\vec{H}, \emptyset, V$);
- 3 **Procedure** BicliqueListing($\vec{H}, R_U, C_V$)
- 4 **if** $|R_U| = \alpha$ **then**
- 5 **foreach** $R_V \subseteq C_V, |R_V| = \beta$ **do** Output (R_U, R_V) as an (α, β) -biclique;
- 6 **return**;
- 7 **foreach** $u \in \vec{H}$ **do**
- 8 $C'_V \leftarrow C_V \cap \text{Nei}_{\vec{G}}(u)$;
- 9 **if** $|C'_V| < \beta$ **or** $\text{Nei}_{\vec{H}}(u) < \alpha - |R_U| - 1$ **then continue**;
- 10 Let $\vec{H}'$ be induced subgraph of $\vec{H}$ on $\text{Nei}_{\vec{H}}(u)$;
- 11 BicliqueListing($\vec{H}', R_U \cup \{u\}, C'_V$);

(U, V, E) , the core idea is to enumerate combinations of vertices in the U set while maintaining their common neighbors in the V set. It also constructs an ordinary directed graph, called the 2-hop graph, on the vertices of the U set to accelerate the enumeration process.

Algorithm 2 presents the basic framework. It first constructs the 2-hop graph $\vec{H}$ with U as its vertex set (Line 1), where an edge (u, v) exists if $u < v$ and u, v share at least one common neighbor in set V of the bipartite graph. In Line 4, if R_U already contains α vertices, it is sufficient to enumerate all combinations of β vertices from C_V to obtain the (α, β) -bicliques. In Line 7, the algorithm enumerates vertices u in $\vec{H}$, then in Line 8 updates the candidate set C'_V in V , and proceeds with the next-level enumeration in Line 11.

C. Union-Find Set [21]

The Union-Find Set (Disjoint-Set) [21] maintains a collection of disjoint sets supporting two operations: **Find**(i) returns the representative of the set containing i , and **Union**(i, j) merges the sets containing i and j . Both operations run in constant amortized time, with $O(n)$ space for n elements.

III. EXISTING SOLUTION

The existing BCPC detection method was proposed in [7], implemented in [15] as a toolkit, and implemented in [13], [14] to build indexes for personalized search problems. To the best of our knowledge, no other work has primarily focused on the algorithmic implementation and improvement of BCPC detection. The basic idea is based on *maximal biclique adjacency graph* (MBAG). This method first enumerates all maximal bicliques and then uses them as vertices to construct a maximal biclique adjacency graph. In this graph, there is an edge between two vertices if and only if the overlapping part of the corresponding two maximal bicliques contains an (α, β) -biclique (satisfying Definition 2). The connected components in the adjacency graph correspond one-to-one with the (α, β) -BCPC. In this method, the (α, β) -biclique connectivity of maximal bicliques in (α, β) -BCPC are maintained by the classic union-find set [21].

Algorithm 3 is the existing MBAG-based algorithm identifying all (α, β) -BCPCs in a bipartite graph $G = (U, V, E)$. It starts by enumerating all maximal bicliques in G and filters them based on size constraints (Lines 1-2). The reason is that

Algorithm 3: Existing MBAG-based solution

Input: Bipartite graph $G = (U, V, E)$; Two integers α, β
Output: All BCPC in G

- 1 $\mathcal{B} \leftarrow$ set of all maximal bicliques;
- 2 $\mathcal{B} \leftarrow \{B \in \mathcal{B}, |B.X| \geq \alpha, |B.Y| \geq \beta\}$;
- 3 $\text{Nei}(B) \leftarrow$ set of maximal bicliques that share vertices with B ;
- 4 $UF \leftarrow$ an initial union-find set, i.e., $UF.\text{find}(B) = B$ for any input B ;
- 5 **for** $B \in \mathcal{B}$ **do**
- 6 **for** $B' \in \text{Nei}(B)$ **do**
- 7 **if** $|B.X \cap B'.X| \geq \alpha$ **and** $|B.Y \cap B'.Y| \geq \beta$ **then**
 $UF.\text{union}(B, B')$;
- 8 **return** UF ;

for a maximal biclique B , if $|B.X| < \alpha$ or $|B.Y| < \beta$, then B can not form a BCPC with other maximal bicliques. A union-find set UF is initialized in Line 4. The algorithm then iterates over each maximal biclique B and its neighbors $B' \in \text{Nei}(B)$, checking if $|B.X \cap B'.X| \geq \alpha, |B.Y \cap B'.Y| \geq \beta$ (Lines 5-7). If the conditions are satisfied, B and B' are merged into the same BCPC by $UF.\text{union}(\ast)$ (Line 7). After processing all maximal bicliques, the union-find set UF represents the final BCPCs in the bipartite graph.

The time complexity of Algorithm 3 is $O(2^{n/2}m + |\mathcal{B}|^2n)$ (maximal biclique enumeration and MBAG traversal. $n = |U| + |V|, m = |E|, \mathcal{B}$ is the set of maximal bicliques). The major drawback of this approach is that the size of the MBAG can be extremely large (up to 10^{11} , see Table II). Consequently, traversing such a massive MBAG to derive all BCPCs is clearly impractical.

IV. NOVEL PARTIAL-BCPC BASED SOLUTION

As we discussed in Section III, current approach relying on MBAG (Algorithm 3) faces significant challenges due to the vast scale of MBAG. This immense size leads to inefficient traversal operations. Consequently, it becomes evident that optimizing the performance of MBAG-based method hinges on effectively reducing the size of the MBAG.

Basic idea. A natural starting point for reducing the size of the MBAG is to leverage the shared prefixes on the maximal biclique enumeration tree. During the enumeration, if a prefix already contains at least α vertices in U and β vertices in V , then all maximal bicliques generated under this prefix share an (α, β) -biclique, so they must belong to the same BCPC. Connections among these maximal bicliques in the MBAG are therefore unnecessary to traverse. We refer to such a set of maximal bicliques as a partial-BCPC. By treating each partial-BCPC as an individual vertex and retaining only the edges between different partial-BCPCs, a logically condensed MBAG with fewer vertices is obtained.

Example 4. Let $\alpha = \beta = 1$. Consider Figure 3, where the subtree rooted at node $nd(B_1, h_2)$ contains three maximal biclique branches: B_1, B_2 , and B_3 . All three maximal bicliques include both $nd(B_1, h_2).R_U$ and $nd(B_1, h_2).R_V$. Therefore, we can regard B_1, B_2 , and B_3 as forming an incomplete BCPC, i.e., a partial-BCPC.

Next, we will formally define partial-BCPC, the maximal biclique enumeration tree and the special prefixes used to compute partial-BCPCs.

Definition 5 (partial-BCPC). Given a bipartite graph G and two integers α, β , a partial-BCPC $\mathcal{P}$ is a set of maximal

bicliques satisfying $\mathcal{P} \subseteq \mathcal{B}$, where $\mathcal{B}$ is an (α, β) -BCPC in G .

The distinction from Definition 4 lies in the fact that Definition 5 does not require the set of maximal bicliques in the partial-BCPC to be maximal. This reflects the characteristic of partial-BCPC as an incomplete BCPC.

Definition 6 (Maximal Biclique Enumeration Tree (MBE Tree)). *Given a bipartite graph $G = (U, V, E)$. The MBE tree $T = (N, \vec{E})$ represents the search space of the maximal biclique enumeration algorithm on G . Here, N denotes the set of nodes in the tree, where each node $nd \in N$ contains a six-tuple $(R_U, R_V, C_U, C_V, X_U, X_V)$, as defined in the basic framework of MBE in Section II-A. The set $\vec{E}$ consists of directed edges of the form (nd_i, nd_j, u) , where $u \in U$ or $u \in V$. An edge $(nd_i, nd_j, u) \in \vec{E}$ indicates that node nd_j is a child of node nd_i , generated by selecting vertex u from $nd_i.C_U$ or $nd_i.C_V$. A node $nd \in N$ is a leaf node if $nd.C_U \cup nd.C_V = \emptyset$. For a leaf node $nd \in N$, if $nd.X_U \cup nd.X_V = \emptyset$, then $(nd.R_U, nd.R_V)$ is a maximal biclique.*

For simplicity, we use $nd.RC X_U$ and $nd.RC X_V$ to denote $nd.R_U \cup nd.C_U \cup nd.X_U$ and $nd.R_V \cup nd.C_V \cup nd.X_V$, respectively.

Definition 7 ((α, β) -Prefix and (α, β) -Node). *Given a bipartite graph $G = (U, V, E)$ and its MBE tree $T = (N, \vec{E})$, let $nd_1 \in N$ be the root node of T , $PF(nd_1) = (nd_1, nd_2, \dots, nd_l)$ is the path from nd_1 to nd_l , i.e., $\forall i = 1, 2, \dots, l-1, (nd_i, nd_{i+1}, *) \in \vec{E}$. $PF(nd_l)$ is an (α, β) -prefix if and only if $|nd_l.R_U| = \alpha, |nd_{l-1}.R_U| < \alpha, |nd_l.R_V| \geq \beta$ or $|nd_l.R_V| = \beta, |nd_{l-1}.R_V| < \beta, |nd_l.R_U| \geq \alpha$. In this case, nd_l is an (α, β) -node.*

In short, an (α, β) -node nd is the first node that satisfies $|nd.R_U| \geq \alpha$ and $|nd.R_V| \geq \beta$ in all branches where nd resides. For all its ancestor nodes nd_i , it is not possible to simultaneously satisfy $|nd_i.R_U| \geq \alpha$ and $|nd_i.R_V| \geq \beta$.

Focusing on (α, β) -nodes. Ancestors of an (α, β) -node are ignored as they fail to satisfy $|R_U| \geq \alpha$ and $|R_V| \geq \beta$. Descendants of an (α, β) -node, while satisfying $|R_U| \geq \alpha$ and $|R_V| \geq \beta$, are redundant since any maximal biclique branches sharing the prefix corresponding to the descendants must also share the (α, β) -prefix corresponding to the (α, β) -node. Thus, focusing solely on (α, β) -nodes enables efficient computation of partial-BCPCs.

A. Basic Framework: Within-Subtree Partial-BCPC

This subsection presents the basic framework of our partial-BCPC based solution, which directly implements the basic idea above by merging only the maximal bicliques located within the subtree rooted at each (α, β) -node. The framework consists of two fundamental steps: (1) Computing partial-BCPCs using the MBE tree; (2) Traversing the reduced MBAG based on the partial-BCPCs to derive the final BCPCs. When computing partial-BCPCs, the framework connects the maximal bicliques within the subtree rooted at (α, β) -node in the MBE tree using a union-find set.

Algorithm 4: Basic within-subtree partial-BCPC based solution

Input: Bipartite graph $G = (U, V, E)$; Two integers α, β
Output: All BCPC in G

```

1 Let  $T = (N, \vec{E})$  be the MBE tree of  $G$ ,  $\mathcal{B}$  be the set of maximal bicliques in  $G$ ;
2  $Nei(B) \leftarrow$  set of maximal bicliques that share vertices with  $B$ ;
3 Let  $UF$  be an initial union-find set;
4 Let  $\mathcal{P} \leftarrow \emptyset$ ;
5 /* compute partial-BCPCs */
6 for  $(\alpha, \beta)$ -node  $nd$  in  $T$  do
7   ProcessNode( $nd$ );
8   for  $B \in \mathcal{P}$  do  $UF.union(B, \mathcal{P}[0])$ ;
9    $\mathcal{P} \leftarrow \emptyset$ ;
10 /* traverse MBAG with partial-BCPCs */
11 for  $B \in \mathcal{B}$  do
12   for  $B' \in Nei(B)$  and  $B, B'$  are in different partial-BCPCs do
13     if  $|B.X \cap B'.X| \geq \alpha, |B.Y \cap B'.Y| \geq \beta$  then
14        $UF.union(B, B')$ ;
15 return  $UF$ ;
16 /* recursively traverse the MBE tree to gather all maximal bicliques under node */
17 Procedure ProcessNode( $node$ )
18 for  $(node, node', u) \in \vec{E}$  do
19   if  $node'.C_U \cup node'.C_V = \emptyset$  then
20     if  $node'.X_U \cup node'.X_V = \emptyset$  then
21        $\mathcal{P} \leftarrow \mathcal{P} \cup \{(node'.R_U, node'.R_V)\}$ ;
22       /*  $(node'.R_U, node'.R_V)$  is a maximal biclique */
23     else ProcessNode( $node'$ );

```

Implementation. Algorithm 4 presents the basic framework. It begins by maximal biclique enumeration and extracting MBE tree $T = (N, \vec{E})$ (Line 1). It then initializes the neighbor list of maximal bicliques and a union-find set UF to manage the partial-BCPCs and BCPCs (Lines 2-3). In Lines 5-8, the algorithm iterates over each (α, β) -node nd in T (Line 5), invoking the ProcessNode function to gather all maximal bicliques under subtree rooted at nd into $\mathcal{P}$ (Line 6). Subsequently, it merges all maximal bicliques in $\mathcal{P}$ into the same partial-BCPC (Line 7). In Lines 9-11, the algorithm traverses a smaller MBAG based on the partial-BCPCs. Specifically, for each maximal biclique B , the algorithm only traverses maximal bicliques B' that reside in different partial-BCPCs from B (Line 10). The algorithm finally returns UF , which represents all detected BCPCs.

The storage of the MBE tree. In practice, we do not store all branches of the MBE tree. We first define two types of nodes in the MBE tree.

Definition 8 (Real and Virtual Node). *Given a graph G and its MBE tree $T = (N, \vec{E})$, for any node $nd \in N$, if nd is a leaf node and satisfies $nd.X_U = nd.X_V = \emptyset$, then nd is a real node—in this case, $(nd.R_U, nd.R_V)$ is a maximal biclique (see Definition 6). Otherwise, if nd has at least one descendant that is a real node, it is also a real node. If neither of these conditions is satisfied, then nd is a virtual node.*

The ProcessNode procedure in Algorithm 4 only needs to identify the branches composed of real nodes (Lines 15-17), i.e., branches containing maximal bicliques. Thus, in practice, we only store real nodes in the MBE tree.

Analysis. We first analyze the correctness of Algorithm 4, then examine the relationship between the number of partial-BCPCs obtained by Algorithm 4 and the numbers of maximal bicliques as well as BCPCs, and finally analyze the complexity

of Algorithm 4.

Theorem 1. *Algorithm 4 correctly computes all BCPCs.*

Proof. We first prove that the partial-BCPCs are all correct. The procedure `ProcessNode` collects all maximal bicliques within the subtree rooted at the given *node* into $\mathcal{P}$. Since the parameter is always an (α, β) -node, the maximal bicliques collected in each execution of the procedure must share an (α, β) -biclique. Therefore, these bicliques correctly form a partial-BCPC. Then we prove that the final BCPCs are correct. In Lines 9-11, any two maximal bicliques that belong to different partial-BCPCs but satisfy the adjacency condition are connected via the union-find set UF . As a result, each disjoint set maintained by UF corresponds to a valid BCPC. $\square$

Theorem 2. *For any two (α, β) -nodes nd_1, nd_2 , nd_1 can not be an ancestor or a descendant of nd_2 .*

Proof. Based on Definition 7, nd_1 is the first node that satisfies $|nd_1.R_U| \geq \alpha$ and $|nd_1.R_V| \geq \beta$ in all branches where nd_1 resides. The same applies to nd_2 . Thus, nd_1, nd_2 are in two different branches, and nd_1 can not be an ancestor or a descendant of nd_2 . $\square$

This non-ancestor property is useful for establishing the counting relationship in the next theorem.

Theorem 3. *Given a bipartite graph G , two integers α, β and MBE tree $T = (N, \vec{E})$, let $n_{biclique}$ be the number of maximal bicliques, $ND \leftarrow \{nd | nd \in N, nd \text{ is an } (\alpha, \beta)\text{-node, } nd \text{ is a real node}\}$, $pbcpc$ be the number of partial-BCPCs computed by Algorithm 4, $bcpc$ be the number of BCPCs, the following inequality holds: $n_{biclique} \geq |ND| = pbcpc \geq bcpc$.*

Proof. Based on Theorem 2, any two (α, β) -nodes in ND can not be in the same branch, and nodes in ND are all real nodes, thus, each node in ND has at least one maximal biclique in its subtree. As a result, $n_{biclique} \geq |ND|$. Since Algorithm 4 connects all maximal bicliques in the subtree of each (α, β) -node in ND into a partial-BCPC, and the maximal bicliques from different subtrees of (α, β) -nodes are distinct, thus, each (α, β) -node in ND corresponds to one partial-BCPC ($|ND| = pbcpc$). Partial-BCPCs are incomplete BCPCs according to Definition 5, thus, $pbcpc \geq bcpc$. $\square$

Theorem 4. *The time and space complexity of Algorithm 4 are $O(2^{n/2}m + (1 - \frac{1}{pbcpc})n_{biclique}^2n)$ and $O(n_{biclique}n)$ respectively, where $n = |U| + |V|$, $m = |E|$, $pbcpc$ is the number of partial-BCPCs, $n_{biclique}$ is the number of maximal bicliques.*

Proof. The time complexity of maximal biclique enumeration is $O(2^{n/2}m)$ [16]. Compared to maximal biclique enumeration, Algorithm 4 introduces additional operations: traversing (α, β) -nodes and collecting maximal biclique branches from the subtrees rooted at these nodes. Since (α, β) -nodes do not reside in the same branch (Theorem 2), the time complexity of these additional operations does not exceed $O(2^{n/2}m)$. In Lines 9-11, in the worst case, Algorithm 4 only skips the inner mutual visits of maximal bicliques in each initial partial-BCPC. Let $x = n_{biclique}$, $y = pbcpc$ and $p_1, p_2, \dots, p_y$ be the number of maximal bicliques in each partial-BCPC, the time complexity of Lines 9-11 is $O(x^2n - (p_1^2 + p_2^2 + \dots + p_y^2)n) =$

$O(x^2n - (\frac{x}{y})^2yn)$ (based on Cauchy-Schwarz inequality), which can be written as $O((1 - \frac{1}{pbcpc})n_{biclique}^2n)$. As a result, the time complexity of Algorithm 4 is $O(2^{n/2}m + (1 - \frac{1}{pbcpc})n_{biclique}^2n)$.

As for the space complexity, the main space consumption in Algorithm 4 comes from the MBE tree and maximal bicliques. Since only the real nodes in the MBE tree are stored, the number of branches in the MBE tree is equal to the number of maximal bicliques, thus, the space complexity of Algorithm 4 is $O(n_{biclique}n)$. $\square$

Discussion. $n_{biclique} \geq pbcpc$ in Theorem 3 forms the basis for reducing the MBAG using partial-BCPC. This is because maximal bicliques within the same partial-BCPC are treated as a single vertex in the MBAG, which leads to an MBAG with formally fewer vertices. This can also be seen from the time complexity in Theorem 4: if $pbcpc$ is replaced with $n_{biclique}$, the result is exactly the complexity of the MBAG-based method in Section III ($O(2^{n/2}m + n_{biclique}^2n)$).

B. Key Optimization: Expanded Cross-Subtree Partial-BCPC

The basic framework in Section IV-A only merges the maximal bicliques located within the subtree of each (α, β) -node. In fact, an (α, β) -node nd may correspond to many more maximal bicliques than those in its own subtree, and missing them leaves redundancy in the resulting MBAG. We therefore introduce a *cross-subtree* expansion strategy that safely identifies and merges the maximal bicliques sharing $(nd.R_U, nd.R_V)$ but scattered across different subtrees into the same partial-BCPC.

For an (α, β) -node nd , not all maximal bicliques sharing $nd.R_U$ and $nd.R_V$ are located within the subtree rooted at nd . For example, in Figure 4 (a), nd is an (α, β) -node, but B_1 and B_2 are not in the subtree rooted at nd , even though both B_1 and B_2 share $nd.R_U$ and $nd.R_V$. This phenomenon occurs because the vertex ordering and pivoting strategy of maximal biclique enumeration place some vertices in B_1 and B_2 (e.g., u_3) early in the candidate set C_U of the root node, so B_1 and B_2 are enumerated earlier than B_4, B_5 and end up in different subtrees of the MBE tree from nd . Yet, by the definition of partial-BCPC, all such cross-subtree maximal bicliques should still be merged into the same partial-BCPC as those in the subtree of nd . Identifying and merging these cross-subtree maximal bicliques yields a substantially larger and more powerful partial-BCPC, which we refer to as an *cross-subtree* partial-BCPC.

Challenges in identifying cross-subtree maximal bicliques.

A naive way to identify cross-subtree maximal bicliques sharing $nd.R_U$ and $nd.R_V$ is to perform a blind traversal of the entire MBE tree starting from its root. However, maximal biclique enumeration is inherently exponential: the MBE tree may contain up to $O(2^{n/2})$ maximal bicliques in the worst case [16]. Without effective pruning, locating such cross-subtree maximal bicliques degenerates into an exhaustive enumeration over this exponential search space.

Targeted search via local information. We address the above challenge by leveraging the local R, C, X sets maintained at the (α, β) -node nd . These sets provide a simple necessary




Fig. 4. Examples of cross-subtree maximal bicliques. $\alpha = 1, \beta = 1$. B_1, B_2, B_4, B_5 (enclosed in boxes) are maximal bicliques sharing $nd.R_U, nd.R_V$

condition for deciding whether another branch can generate maximal bicliques sharing the prefix $(nd.R_U, nd.R_V)$.

Theorem 5. Let nd be an (α, β) -node in the MBE tree, and let nd' be any other node. If some maximal biclique generated in the subtree rooted at nd' contains $nd.R_U$ and $nd.R_V$, then $nd'.R_U \subseteq nd.RC X_U$ and $nd'.R_V \subseteq nd.RC X_V$.

Proof. Let B be a maximal biclique generated in the subtree rooted at nd' and containing $nd.R_U$ and $nd.R_V$. Since B is generated below nd' , the current biclique of nd' must be contained in B , i.e., $nd'.R_U \subseteq B.X$ and $nd'.R_V \subseteq B.Y$. By the properties of maximal biclique enumeration, any maximal biclique containing $nd.R_U$ and $nd.R_V$ has all its vertices confined within $nd.RC X_U$ and $nd.RC X_V$ [16]. Thus, $nd'.R_U \subseteq nd.RC X_U$ and $nd'.R_V \subseteq nd.RC X_V$. $\square$

By Theorem 5, any node failing this condition can be safely skipped. Thus, the global traversal is reduced to a targeted search over branches that may contain the prefix $(nd.R_U, nd.R_V)$; the maximal bicliques found there are merged into the same cross-subtree partial-BCPC.

Stop-Label. The targeted search above may still traverse many branches down to leaf nodes. To reduce this overhead, we assign *stop-labels* to real (α, β) -nodes and their immediate real child nodes; each label is an arbitrary maximal biclique in a partial-BCPC and represents the partial-BCPC containing all maximal bicliques in the corresponding subtree. Within the subtree of an (α, β) -node, this is straightforward: we only collect the stop-labels of its immediate real child nodes. In beyond-subtree search, the search can also stop when it reaches

the stop-label of another (α, β) -node; the following theorem gives the reason why that outside subtree can be collected as a whole.

Theorem 6. Let nd be an (α, β) -node in the MBE tree. For any other (α, β) -node nd' that satisfies $nd'.R_U \subseteq nd.RC X_U$ and $nd'.R_V \subseteq nd.RC X_V$, every maximal biclique generated in the subtree rooted at nd' belongs to the same BCPC as the maximal bicliques in the subtree rooted at nd .

Proof. By the definition of an (α, β) -node, the sets $(nd.R_U, nd.R_V)$ form a complete bipartite subgraph. Since $nd.RC X_U$ and $nd.RC X_V$ respectively consist of all vertices in G that maintain a complete bipartite relationship with $nd.R_V$ and $nd.R_U$, the condition $nd'.R_U \subseteq nd.RC X_U$ and $nd'.R_V \subseteq nd.RC X_V$ implies that every vertex in $nd'.R_U$ is adjacent to every vertex in $nd.R_V$, and vice versa. Consequently, the union $B^* = (nd.R_U \cup nd'.R_U, nd.R_V \cup nd'.R_V)$ is a biclique. By the property of maximal biclique enumeration, there must exist at least one maximal biclique B_{bridge} in G such that $B^* \subseteq B_{bridge}$. Let B_{nd} be any maximal biclique generated from the subtree of nd , and $B_{nd'}$ be any maximal biclique from the subtree of nd' . B_{nd} and B_{bridge} are connected in the MBAG because they share $(nd.R_U, nd.R_V)$. Similarly, $B_{nd'}$ and B_{bridge} are connected as they share $(nd'.R_U, nd'.R_V)$. Due to this shared adjacency with B_{bridge} , all such maximal bicliques belong to the same BCPC. $\square$

We implement stop-labels as follows:

Within-Subtree Stop-Labels. For real child nodes of an (α, β) -node, stop-labels are computed bottom-up. If a child is a leaf, its stop-label is the maximal biclique represented by the leaf; otherwise, the stop-labels of its real child nodes are merged into one partial-BCPC, and any maximal biclique in it becomes the current child's stop-label.

Beyond-Subtree Stop-Labels. If the search for nd reaches another (α, β) -node nd' satisfying Theorem 6 and already having a stop-label, it stops at nd' and directly collects the represented partial-BCPC. All collected stop-labels are finally merged into one partial-BCPC, whose representative maximal biclique becomes the stop-label of nd .

Example 5. Let $\alpha = 1, \beta = 1$, when processing nd in Figure 4 (b):

During the search of subtree rooted at nd , the process terminates upon encountering the stop-label of child node of nd . This stop-label is B_4 , with B_5 belonging to the same partial-BCPC as B_4 .

During the search beyond the subtree rooted at nd , the search process terminates upon detecting the stop-label of nd' .

Finally, the two stop-labels (B_1 and B_4) are connected; concurrently, the partial-BCPCs they belong to are merged into a larger partial-BCPC (including B_1, B_2, B_4, B_5).

Implementation. Algorithm 5 integrates the above strategies into the basic framework via a post-order traversal (Lines 4, 7). For each (α, β) -node and its descendants, `ProcessNode+` (Lines 10, 12) identifies relevant maximal bicliques or existing stop-labels and collects them in $\mathcal{P}$. These are then unified into a single partial-BCPC, with the first element designated as the node's stop-label (Lines 13-15). Specifically, when

Algorithm 5: Expanded cross-subtree partial-BCPC based solution

```

Input: Bipartite graph  $G = (U, V, E)$ ; Two integers  $\alpha, \beta$ 
Output: All BCPC in  $G$ 
1 Lines 1-3 of Algorithm 4;
2 Let  $\mathcal{P} \leftarrow \emptyset$ ,  $root \leftarrow$  root node of  $T$ ;
3 Let  $ST(nd)$ ,  $nd \in N$  be the stop-label of  $nd$ ;
  /* compute partial-BCPCs */
4 Postorder( $root$ );
  /* traverse MBAG with partial-BCPCs */
5 Lines 9-11 of Algorithm 4;
6 return  $UF$ ;

7 Procedure Postorder( $node$ )
8   foreach  $child$  of  $node$ ,  $node'$ , in creation order of  $node'$  do Postorder( $node'$ );
9   if  $node$  is an  $(\alpha, \beta)$ -node then
10     ProcessNode+
      ( $root$ ,  $node.RCX_U \cup node.RCX_V$ ,  $node.R_U \cup node.R_V$ );
11   else if  $node$  is a real node and a descendant of an  $(\alpha, \beta)$ -node then
12     ProcessNode+( $node$ );
13   if  $\mathcal{P}$  is not empty then
14     for  $B \in \mathcal{P}$  do  $UF.union(B, \mathcal{P}[0])$ ;
15      $ST(node) \leftarrow \mathcal{P}[0]$ ,  $\mathcal{P} \leftarrow \emptyset$ ;

16 Procedure ProcessNode+ ( $node$ ,  $RCX$ ,  $R$ )
17   foreach  $child$  of  $node$ ,  $node'$ , in creation order of  $node'$  do
18     Let ( $node$ ,  $node'$ ,  $u$ )  $\in \vec{E}$ ;
19     if  $node'$  is a real node then
20       /* targeted search: only maximal bicliques
21        selecting vertices in  $RCX$  may contain the
22        target prefix  $R$  */
23       if  $u \in RCX$  then
24         if  $ST(node') = None$  then ProcessNode+
25           ( $node'$ ,  $RCX$ ,  $R$ );
26         else  $\mathcal{P} \leftarrow \mathcal{P} \cup \{ST(node')\}$ ;
27       if  $u \in R$  then break;

28 Procedure ProcessNode+ ( $node$ )
29   if  $node$  is a leaf node then  $\mathcal{P} \leftarrow \mathcal{P} \cup \{(node.R_U, node.R_V)\}$ ;
30   else
31     foreach  $node'$ ,  $node'$  is a child of  $node$  and a real node do
32        $\mathcal{P} \leftarrow \mathcal{P} \cup \{ST(node')\}$ 

```

processing an (α, β) -node (Line 9), ProcessNode+ traverses all its real nodes (Lines 17-19) and initiates targeted recursive search (Line 20) that terminates upon encountering a pre-assigned stop-label (Line 22). Following the MBE enumeration principle, after the branch selecting vertex u (Line 18) has been processed, no subsequent sibling branch can contain u ; hence, if $u \in R$, those branches cannot contain the whole target set R and can be skipped (Line 23). For descendant nodes (Line 11), ProcessNode+ either adds the represented maximal biclique (if a leaf, Line 25) or merges the stop-labels of all child nodes into $\mathcal{P}$ (Lines 27-28). Finally, the algorithm obtains the BCPCs by traversing the MBAG based on these partial-BCPCs (Line 5).

Analysis. We now proceed to analyze the correctness of Algorithm 5. The key lies in demonstrating that the partial-BCPCs computed in Postorder (Line 4) are correct.

Theorem 7. For each (α, β) -node processed in Algorithm 5 (Line 9), ProcessNode+ connects all maximal bicliques in the bipartite graph that share $(node.R_U, node.R_V)$ into the same partial-BCPC.

Proof. Consider a sequence of (α, β) -nodes $(nd_1, nd_2, \dots, nd_n)$, which follow the processing order in Postorder (i.e., post-order traversal). We first prove that during the processing of nd_1 , ProcessNode+ connects all maximal bicliques that share $(nd_1.R_U, nd_1.R_V)$ into the

same partial-BCPC. Since nd_1 is the first (α, β) -node, it must also be in the first enumerated maximal biclique branch (if nd_1 appears in an earlier branch, then that earlier or an even earlier branch is likely the first one containing a maximal biclique; if nd_1 appears in a later branch, then it can not be the first (α, β) -node). Therefore, all maximal biclique branches that contain $(nd_1.R_U, nd_1.R_V)$ must lie in the subtree rooted at nd_1 , and can not appear in any subtree processed before nd_1 .

Using mathematical induction, assume that ProcessNode+ connects all maximal bicliques that share $(nd_i.R_U, nd_i.R_V)$ when processing $nd_1, nd_2, \dots, nd_i$. Next, we prove the case for $i + 1$. For nd_{i+1} , let $R_U = nd_{i+1}.R_U$ and $R_V = nd_{i+1}.R_V$. According to the MBE process [16], the branches containing maximal bicliques with R_U and R_V are either generated before nd_{i+1} , or located in the subtree rooted at nd_{i+1} . For the maximal biclique branches generated before nd_{i+1} , their corresponding (α, β) -node nd_j must satisfy $j < i + 1$, so we only need to collect $ST(nd_j)$, as it represents the partial-BCPCs of those maximal bicliques. For the maximal bicliques in the subtree of nd_{i+1} , we only need to collect the stop-labels of the child nodes of nd_{i+1} , because the maximal bicliques under each child node have already been connected into partial-BCPCs (since the traversal is post-order). Therefore, the case for nd_{i+1} is proved. $\square$

Theorem 8. Algorithm 5 correctly computes all BCPCs.

Proof. The proof of this theorem is similar to that of Theorem 1. Since Theorem 7 proves that the partial-BCPCs in Algorithm 5 are obtained by connecting maximal bicliques that share the R_U and R_V of (α, β) -nodes, these partial-BCPCs are correct. Moreover, the process of computing the final BCPCs in Algorithm 5 (Line 5) is the same as in Algorithm 4, thus, Algorithm 5 can correctly compute the BCPCs. $\square$

Theorem 9. Given a bipartite graph G , two integers α, β , let $pbcpc$ be the number of within-subtree partial-BCPCs computed by Algorithm 4, $pbcpc+$ be the number of cross-subtree partial-BCPCs computed by Algorithm 5, $bpcpc$ be the number of BCPCs, the following inequality holds: $pbcpc \geq pbcpc+ \geq bpcpc$.

Proof. We only need to prove $pbcpc \geq pbcpc+$. When processing an (α, β) -node, Algorithm 5 not only connects maximal bicliques in the subtree rooted at that node, but also further connects maximal bicliques that were enumerated in earlier branches and share R_U, R_V of the (α, β) -node. Moreover, the ProcessNode+ procedure in Algorithm 5 also processes (α, β) -nodes that belong to virtual nodes. Therefore, $pbcpc \geq pbcpc+$. $\square$

Theorem 10. The time and space complexity of Algorithm 5 are $O(2^{n/2}mn^d + (1 - \frac{1}{pbcpc+})n_{biclique}^2n)$ and $O(n_{biclique}n)$ respectively, where $n = |U| + |V|$, $m = |E|$, d is the maximum depth of real (α, β) -nodes (Definition 8), $pbcpc+$ is the number of partial-BCPCs, $n_{biclique}$ is the number of maximal bicliques.

Proof. The time complexity of maximal biclique enumeration is $O(2^{n/2}m)$ [16]. Unlike Algorithm 4, Algorithm 5 handles not only (α, β) -nodes (Lines 9-10) but also their child nodes

Similar to Algorithm 4, the space complexity of Algorithm 5 is also $n_{biclique}n$. $\square$

V. NOVEL (α, β) -BICLIQUE BASED SOLUTION

A. The Basic Framework

Implementation. Algorithm 6 presents the basic framework of the (α, β) -biclique based solution. It is built upon the (α, β) -biclique enumeration framework proposed in [17], where the core idea is to enumerate combinations of vertices in the U set while maintaining their common neighbors in the V set. This framework constructs an ordinary directed graph, called the *2-hop graph*, on the vertices of the U set to accelerate the enumeration process.

Algorithm 6: (α, β) -Biclique based solution

```

1 Let  $UF$  be an initial union-find set;
2  $\mathcal{B} \leftarrow$  set of all maximal bicliques in  $G$ ;
3  $\mathcal{B} \leftarrow \{B | B \in \mathcal{B}, |B.X| \geq \alpha, |B.Y| \geq \beta\}$ ;
4 Construct 2-hop graph  $\vec{H} = (U, \vec{E}_h)$  on vertex set  $U$  [17];
    $/* \vec{H} = (U, \vec{E}_h):$  a directed graph,  $(u, v) \in \vec{E}_h$  indicates
   that  $u < v$ ,  $u, v$  have common neighbors in  $V$   $*/$ 
5 BicliqueListing( $\vec{H}, \emptyset, V$ );
6 return  $UF$ ;

7 Procedure BicliqueListing( $\vec{H}, R_U, C_V$ )
8 if  $|R_U| = \alpha$  then
9   foreach  $R_V \subseteq C_V, |R_V| = \beta$  do Connect maximal bicliques sharing
    $(R_U, R_V)$  in  $\mathcal{B}$  with  $UF.union(*)$ ;
10  return;

11 foreach  $u \in \vec{H}$  do
12    $C'_V \leftarrow C_V \cap Nei_G(u)$ ;
13   if  $|C'_V| < \beta$  or  $Nei_{\vec{H}}(u) < \alpha - |R_U| - 1$  then continue;
14   Let  $\vec{H}'$  be induced subgraph of  $\vec{H}$  on  $Nei_{\vec{H}}(u)$ ;
15   BicliqueListing( $\vec{H}', R_U \cup \{u\}, C'_V$ );

```

Pruning with Maximal Bicliques. The key to addressing the above challenge is to reduce the search space in the enumeration process of (α, β) -bicliques — that is, to prune branches whose associated maximal bicliques already belong to the same BCPC.

Algorithm 7: (α, β) -Biclique based solution with maximal biclique

Input: Bipartite graph $G = (U, V, E)$; Two integers α, β
Output: All BCPC in G

```

1 Let  $UF$  be an initial union-find set;
2  $B \leftarrow$  set of all maximal bicliques in  $G$ ;
3  $B \leftarrow \{B | B \in B, |B.X| \geq \alpha, |B.Y| \geq \beta\}$ ;
4 Construct 2-hop graph  $\vec{H} = (U, \vec{E}_h)$  on vertex set  $U$  [17];
   /*  $\vec{H} = (U, \vec{E}_h)$ : a direct graph,  $(u, v) \in \vec{E}_h$  indicates
   that  $u < v$ ,  $u, v$  have common neighbors in  $V$  */
5 BicliqueListing+ ( $\vec{H}, \emptyset, V$ );
6 return  $UF$ ;

7 Procedure BicliqueListing+ ( $\vec{H}, R_U, C_V$ )
8  $RB \leftarrow \{B | R_U \subseteq B.X, B \in B\}$ ; /* maximal bicliques
   sharing  $R_U$  */
9 if  $|\{UF.find(B) | B \in RB\}| \leq 1$  then return; /* there is only 1
   or 0  $UF$  disjoint set in  $RB$  */
10 if  $|R_U| = \alpha$  then
11   foreach  $R_V \subseteq C_V, |R_V| = \beta$  do
12      $RB \leftarrow \{B | R_V \subseteq B.Y, B \in RB\}$ ;
13     if  $|\{UF.find(B) | B \in RB\}| \leq 1$  then continue; /* there
       is only 1 or 0  $UF$  disjoint set in  $RB$  */
14     Connect maximal bicliques in  $RB$  with  $UF.union(*)$ ;
15   return;
16 foreach  $u \in \vec{H}$  do
17    $C'_V \leftarrow C_V \cap Nei_G(u)$ ;
18   if  $|C'_V| < \beta$  or  $Nei_{\vec{H}}(u) < \alpha - |R_U| - 1$  then continue;
19   Let  $\vec{H}'$  be induced subgraph of  $\vec{H}$  on  $Nei_{\vec{H}}(u)$ ;
20   BicliqueListing+ ( $\vec{H}', R_U \cup \{u\}, C'_V$ );

```

Theorem 13. Given a bipartite graph G , for a biclique B , if the number of maximal bicliques sharing it is 0, or all those maximal bicliques are in the same BCPC, then for any other biclique B' such that $B \subseteq B'$, the number of maximal bicliques sharing B' is also 0, or they are in the same BCPC.

Proof. It is easy to see that any maximal biclique containing B' must also contain B , so the set of maximal bicliques containing B' is a subset of those containing B . Therefore, Theorem 13 can be easily proved. $\square$

Theorem 13 reveals the core idea of pruning the (α, β) -biclique enumeration space: during the enumeration process, for each intermediate small biclique generated, we can maintain the set of maximal bicliques containing it. If this set is empty or all the maximal bicliques in the set already belong to the same BCPC, then the current enumeration branch can be safely pruned.

Implementation. Algorithm 7 presents the implementation of pruning the (α, β) -biclique enumeration process using maximal bicliques. The difference between Algorithm 7 and Algorithm 6 lies in the BicliqueListing+ function. Specifically, in Line 8, BicliqueListing+ collects maximal bicliques that share $(R_U, \emptyset)$. In Line 9, if the maximal bicliques in RB belong to only one or zero union-find sets, then the entire enumeration branch can be pruned. In Lines 12–13, the algorithm further checks whether the maximal bicliques in RB that share R_V belong to more than one union-find set. If so, it connects them using $UF.union(*)$.

Analysis. Now we proceed to analyze the correctness and complexity of Algorithm 7.

Theorem 14. Algorithm 7 correctly computes all BCPCs.

Proof. Compared with Algorithm 6, Algorithm 7 prunes the branches that can not connect multiple UF sets based on

Algorithm 8: (α, β) -Biclique based solution with partial-BCPC

Input: Bipartite graph $G = (U, V, E)$; Two integers α, β
Output: All BCPC in G

```

1 Lines 1-3 of Algorithm 5; /* cross-subtree partial-BCPC
   computation */
2 Lines 2-5 of Algorithm 7; /* connect maximal bicliques by
    $(\alpha, \beta)$ -biclique listing */
3 return  $UF$ ;

```

Theorem 13. Therefore, Algorithm 7 maintains the same correctness as Algorithm 6. $\square$

Theorem 15. The time and space complexity of Algorithm 7 are $O((a(\vec{H}_m)^{\alpha-2}|\vec{E}_{hm}||V_m| + \Delta_m)n_{biclique})$ and $O(n_{biclique}n)$ respectively. $n = |U| + |V|$. Let U_m and V_m be the vertex sets shared by any two maximal bicliques, G_m be the induced subgraph on U_m and V_m , and $\vec{H}_m = (U_m, \vec{E}_{hm})$ be the 2-hop graph based on G_m . $a(\vec{H}_m)$ is the arboricity of $\vec{H}_m$ [17], Δ_m is the number of (α, β) -bicliques in G_m , $n_{biclique}$ is the number of maximal bicliques.

Proof. The time complexity of maximal biclique enumeration is $T_{biclique} = 2^{n/2}|E|$ [16]. Compared with Algorithm 6, Algorithm 7 prunes the branches that can not connect multiple UF sets based on Theorem 13. This implies that vertices initially contained in only one maximal biclique need not be considered during the (α, β) -biclique enumeration process. Therefore, Algorithm 7 can be regarded as performing (α, β) -biclique enumeration on G_m . Considering Lines 8–9 and Lines 12–13 of Algorithm 7, its time complexity is $O(T_{biclique} + (a(\vec{H}_m)^{\alpha-2}|\vec{E}_{hm}||V_m| + \Delta_m)n_{biclique})$. Since $O(n_{biclique}) = O(2^{n/2})$ [22], [16], $O(\Delta_m) = O((U_m + V_m)^{\alpha+\beta})$ [17], the time complexity can be rewritten as $O((a(\vec{H}_m)^{\alpha-2}|\vec{E}_{hm}||V_m| + \Delta_m)n_{biclique})$.

Similarly, the space complexity of Algorithm 7 is still $O(n_{biclique}n)$. $\square$

Pruning with Cross-Subtree Partial-BCPC. Since maximal bicliques in the same partial-BCPC are guaranteed to belong to the same final BCPC (Section IV), they can be pre-merged before (α, β) -biclique enumeration. Compared with pruning by maximal bicliques alone, cross-subtree partial-BCPCs merge more maximal bicliques in advance, making the same pruning condition in Theorem 13 hold earlier and more often.

Implementation. Algorithm 8 is the algorithm that leverages cross-subtree partial-BCPCs, and it consists of two main parts. First, it computes cross-subtree partial-BCPCs (Line 1), which are recorded in the union-find structure (UF). Then, in Line 2, these partial-BCPCs are used to further prune the (α, β) -biclique listing process.

Analysis. Now we proceed to analyze the correctness and complexity of Algorithm 8.

Theorem 16. Algorithm 8 correctly computes all BCPCs.

Theorem 17. The time and space complexity of Algorithm 8 are $O((a(\vec{H}_p)^{\alpha-2}|\vec{E}_{hp}||V_p| + \Delta_p)n_{biclique})$ and $O(n_{biclique}n)$ respectively. $n = |U| + |V|$. Let U_p and V_p be the vertex sets shared by any two partial-BCPCs, G_p be the induced subgraph on U_p and V_p , and $\vec{H}_p = (U_p, \vec{E}_{hp})$ be the 2-hop

TABLE II

DATASETS. THE COLUMNS U , V AND E DESCRIBE THE MEANING OF THE TWO VERTEX SETS AND THE EDGE SET IN EACH BIPARTITE GRAPH; $\bar{d}_U$ IS THE AVERAGE DEGREE OF VERTICES IN U , $\bar{d}_V$ IS THE AVERAGE DEGREE OF VERTICES IN V , N_m IS THE NUMBER OF MAXIMAL BICLIQUES, N_o IS THE NUMBER OF EDGES IN MAXIMAL BICLIQUE ADJACENT GRAPH, $K=10^3$, $M=10^6$, $B=10^9$

Datasets	U	V	E	$ U $	$ V $	$ E $	$\bar{d}_U$	$\bar{d}_V$	N_m	N_o
Youtube	user	group	membership	94,238	30,087	293,360	3.11	9.75	1,769,331	2B
Bookcrossing	user	book	rating	77,802	185,955	433,652	5.57	2.33	155,391	1B
Bibsonomy	tag	publication	assignment	204,674	767,448	2,499,057	12.48	3.32	2,132,826	26B
CiteSeer	author	paper	authorship	105,353	181,395	512,267	4.86	2.82	54,083	962K
Stackoverflow	user	post	favorite	545,195	96,678	1,301,942	2.39	13.47	2,922,148	26B
Twitter	user	hashtag	usage	175,214	530,418	1,890,661	10.79	3.56	5,102,542	13B
Imdb	person	work	association	685,568	186,414	2,715,604	3.96	14.57	1,809,175	1B
Actor2	actor	movie	appearance	303,617	896,302	3,782,463	12.46	4.22	3,761,666	2B
Amazon	user	product	rating	2,146,057	1,230,915	5,743,258	2.68	4.67	5,702,211	108B
DBLP	author	paper	authorship	1,953,085	5,624,219	12,282,059	6.29	2.18	1,281,831	85M

graph based on G_p . $a(\vec{H}_p)$ is the arboricity of $\vec{H}_p$ [17], Δ_p is the number of (α, β) -bicliques in G_p , $n_{biclique}$ is the number of maximal bicliques.

The proofs of Theorem 16 and Theorem 17 are similar to those of Theorem 14 and Theorem 15.

C. Overall Complexity Discussion

A BCPC is built from bicliques, and biclique-related problems such as maximal biclique enumeration [19], [20], [16] and (α, β) -biclique enumeration [17], [18] are well-known to be exponential in the worst case, since the number of such bicliques in a bipartite graph can itself be exponential. Exact BCPC detection therefore inherits an exponential worst-case complexity, and our algorithm is no exception. We note, however, that this exponential bound mostly reflects the theoretical worst case rather than the behavior on real-world bipartite graphs in practice [20], [16], [17].

The focus of this paper is the cross-subtree partial-BCPC, which preserves exactness while reducing redundant computation in two ways: it compresses the MBAG by grouping maximal bicliques that provably lie in the same BCPC (Section IV), and it prunes (α, β) -biclique enumeration by skipping search states whose maximal bicliques are already in a single partial-BCPC (Section V). The experiments in the next section measure these two effects separately, and together they yield the practical improvement of our BCPC detection framework over existing methods.

VI. EXPERIMENTS

A. Experimental Setup

Datasets. We use 10 real-world bipartite graph datasets, whose detailed information is presented in Table II. All the datasets are downloaded from <http://konect.cc/>.

Algorithms. We first implement the state-of-the-art algorithm in [7], [15], [14] as MBAG.

- MBAG (Algorithm 3) is directly based on the maximal biclique adjacency graph (also abbreviated as MBAG). It first enumerates all maximal bicliques, then traverses maximal biclique adjacency graph to compute the BCPC.

We then implement the following partial-BCPC based solution and (α, β) -biclique based solution.

- PBCPC (Algorithm 4) is the basic partial-BCPC based solution. It first enumerates maximal bicliques, and builds the MBE tree. For each (α, β) -node in the MBE tree, it identifies all maximal bicliques enumerated within the subtree rooted at that node and connects them into a

single partial-BCPC. It then traverses the MBAG reduced by the partial-BCPCs to obtain the final BCPC.

- PBCPC+ (Algorithm 5) is an improved version of PBCPC. Unlike PBCPC, it performs cross-subtree search to find maximal bicliques within and beyond the subtree rooted at an (α, β) -node, and uses stop-labels to reduce the search overhead.
- Biclique (Algorithm 6) is the basic solution based on (α, β) -biclique enumeration. Its core idea is to enumerate all (α, β) -bicliques, and for each such biclique, connect all maximal bicliques that contain it in order to obtain the BCPC.
- BicliqueM (Algorithm 7) is an improved version of Biclique. It prunes (α, β) -biclique enumeration branches by checking whether the maximal bicliques associated with the currently enumerated biclique already belong to the same BCPC.
- BicliqueP (Algorithm 8) is also an improved version of Biclique. Unlike BicliqueM, it leverages the method from PBCPC+ to compute cross-subtree partial-BCPCs in advance, thereby identifying some maximal bicliques that already belong to the same BCPC, enabling a stronger pruning effect.

Parameters. There are two parameters used in the experiments: α and β . Their default values are both set to 4, and the possible values they can take are 2, 4, 6, and 8. Recall that α and β control the community cohesion: smaller values yield larger and looser communities, while larger values yield tighter and higher-confidence ones.

All the algorithms are implemented in C++. The experiments are conducted on a system running Ubuntu 20.04.4 LTS with an AMD Ryzen 3900X 2.2GHz CPU and 256GB of memory.

B. Experimental Results

Exp-1: Efficiency of various algorithms. In this experiment, we evaluated the performance of all algorithms across all datasets, and the results are presented in Table III.

Overall, MBAG and PBCPC algorithms perform the worst in terms of runtime across most datasets and α, β . This is mainly because the MBAG algorithm traverses the original maximal biclique adjacency graph (MBAG), which is often very large (see Table II). Although the PBCPC algorithm attempts to reduce the size of the MBAG by computing partial-BCPCs, these partial-BCPCs are not sufficiently comprehensive, leading to limited pruning effectiveness on the MBAG

TABLE III
RUNNING TIME (SECOND) OF ALL ALGORITHMS ON ALL DATASETS. THE BEST AND SUB-BEST RESULTS ARE BOLD AND UNDERLINED RESPECTIVELY.
“-” DENOTES OUT OF TIME (≥ 1 DAYS OR 86,400 SECONDS)

Datasets		Youtube					Bookcrossing				Bibsonomy				Citeseer				Stackoverflow			
α	β	2	4	6	8	2	4	6	8	2	4	6	8	2	4	6	8	2	4	6	8	
2	Algo																					
	MBAG	13,376	11,492	6,072	1,902	35	2	0.9	0.7	9,058	8,096	6,980	5,505	1	0.8	0.7	0.6	32,598	17,967	4,287	844	
	PBCPC	4,651	7,283	6,942	3,207	42	3	1	1	5,173	5,315	5,636	5,188	1	0.8	0.7	0.6	30,051	25,220	8,811	2,177	
	PBCPC+	197	1,914	2,787	1,484	20	3	1	0.8	356	892	1,362	1,736	1	0.9	0.8	0.6	2,564	7,144	3,467	987	
	Biclique	135	358	2,995	-	3	26	-	-	193	-	-	-	1	1	64	6,999	2,038	3,118	-	-	
4	BicliqueM	59	70	43	22	2	1	0.8	0.6	54	52	52	52	1	0.8	0.7	0.6	486	208	119	80	
	BicliqueP	34	43	31	20	3	2	1	1	48	40	39	37	1	0.9	0.8	0.6	339	159	121	103	
	MBAG	13,113	11,358	5,941	1,793	3	0.7	0.5	0.4	8,663	7,755	6,734	5,079	0.1	0.08	0.07	0.06	27,582	14,496	3,308	544	
	PBCPC	7,748	9,568	7,802	3,197	4	0.9	0.6	0.5	5,439	6,427	6,254	5,647	0.1	0.09	0.08	0.06	37,632	25,336	7,496	1,476	
	PBCPC+	1,450	2,964	3,404	1,618	4	0.8	0.6	0.5	1,010	1,474	1,969	2,267	0.1	0.08	0.07	0.06	13,222	9,955	3,857	864	
6	Biclique	3,754	5,445	5,522	2,877	2	1	0.8	0.6	7,946	-	-	-	0.1	0.1	0.1	0.3	1,210	5,446	748	11,595	
	BicliqueM	54	235	140	50	1	0.8	0.6	0.5	106	100	89	82	0.1	0.09	0.07	0.06	306	539	199	56	
	BicliqueP	28	66	49	24	1	0.9	0.6	0.5	53	46	44	39	0.1	0.09	0.08	0.07	190	159	76	34	
	MBAG	10,936	9,409	4,635	1,312	1	0.6	0.5	0.4	7,280	6,859	5,765	4,367	0.03	0.03	0.03	0.03	9,142	4,150	561	45	
	PBCPC	7,736	8,370	5,855	2,080	1	0.7	0.5	0.4	5,219	6,098	6,169	5,605	0.05	0.04	0.04	0.04	12,482	6,893	1,134	89	
8	PBCPC+	2,500	3,622	3,133	1,121	1	0.7	0.5	0.4	1,572	2,050	2,429	2,583	0.04	0.03	0.04	0.03	5,758	3,833	775	76	
	Biclique	-	3,289	34,084	3,352	36	0.9	0.5	0.4	9,143	-	-	-	0.07	0.04	0.04	0.04	-	1,764	2,724	26	
	BicliqueM	44	165	178	35	1	0.6	0.5	0.4	72	138	124	107	0.04	0.04	0.04	0.04	90	131	40	12	
	BicliqueP	24	57	56	16	1	0.7	0.5	0.4	39	72	57	54	0.05	0.04	0.04	0.04	75	72	22	10	
	MBAG	7,236	6,102	2,832	744	1	0.5	0.4	0.3	5,410	5,144	4,299	3,133	0.03	0.03	0.03	0.03	1,377	398	21	7	
	PBCPC	5,585	5,707	3,596	1,241	1	0.6	0.4	0.3	4,829	5,309	5,229	4,508	0.03	0.03	0.03	0.03	1,721	592	32	6	
	PBCPC+	2,327	2,878	1,991	713	1	0.6	0.4	0.3	1,735	2,012	2,291	2,268	0.03	0.03	0.03	0.03	1,110	447	29	7	
	Biclique	-	-	4,728	8,555	2,315	6	0.4	0.3	-	-	-	-	0.09	0.04	0.04	0.04	-	77	20	14	
	BicliqueM	33	117	101	38	1	0.6	0.4	0.3	50	137	178	103	0.03	0.03	0.03	0.03	30	22	9	7	
	BicliqueP	19	45	33	13	1	0.6	0.4	0.3	32	58	85	56	0.04	0.04	0.04	0.03	30	16	9	6	
Datasets		Twitter				Imdb				Actor2				Amazon				DBLP				
α	β	2	4	6	8	2	4	6	8	2	4	6	8	2	4	6	8	2	4	6	8	
2	Algo																					
	MBAG	11,652	5,290	1,817	657	859	706	622	409	7,303	6,650	5,217	3,502	16,282	3,974	580	137	18	12	10	9	
	PBCPC	47,374	24,237	6,055	2,335	323	188	293	418	1,820	2,285	2,377	2,041	20,703	6,687	1,297	342	28	19	15	13	
	PBCPC+	10,580	11,843	3,162	1,337	102	58	90	160	266	524	772	780	2,413	3,067	706	228	23	15	12	10	
	Biclique	685	5,339	4,584	-	55	147	58,284	-	190	875	39,764	-	964	727	-	-	20	176	-	-	
4	BicliqueM	510	777	208	108	35	26	23	21	88	56	48	47	618	157	79	55	17	13	10	9	
	BicliqueP	518	402	152	94	42	33	28	25	79	51	44	39	590	176	108	82	25	19	16	14	
	MBAG	4,339	2,737	1,236	457	762	687	614	403	7,107	6,586	5,178	3,468	7,856	1,720	227	57	7	5	4	3	
	PBCPC	11,436	6,764	3,176	1,482	408	355	446	514	1,622	2,251	2,462	2,125	9,234	2,496	427	100	8	5	4	3	
	PBCPC+	6,282	4,237	2,004	1,001	99	74	121	185	201	516	794	840	4,219	1,508	304	78	8	5	4	3	
6	Biclique	-	12,028	20,544	-	92	1,805	1,335	25,431	626	4,996	29,168	-	534	2,687	474	9,676	12	7	29	1,834	
	BicliqueM	134	129	94	87	22	21	26	21	57	131	51	44	303	161	61	39	8	5	4	3	
	BicliqueP	134	100	69	64	23	18	16	15	47	48	34	31	280	111	61	42	9	5	4	3	
	MBAG	2,184	1,596	891	298	623	600	536	340	6,859	6,376	5,085	3,343	1,576	349	63	32	3	2	2	2	
	PBCPC	3,124	2,404	1,677	690	557	551	605	568	1,988	2,585	2,583	2,270	1,669	499	97	44	3	2	2	2	
8	PBCPC+	2,666	1,469	1,073	506	169	164	204	237	326	641	933	955	1,035	331	71	36	3	2	2	2	
	Biclique	-	-	-	27,608	1,748	835	11,021	10,968	-	22,912	49,668	-	33,212	1,595	525	806	10,632	6	17	30	
	BicliqueM	90	81	68	52	17	17	22	23	51	121	62	52	171	80	44	32	3	2	2	2	
	BicliqueP	95	58	51	41	19	14	14	13	40	43	38	32	213	79	47	33	3	2	2	2	
	MBAG	1,482	1,166	651	210	400	388	340	187	6,381	5,838	4,529	2,958	445	104	41	28	1	1	1	1	
	PBCPC	1,479	1,297	994	369	440	437	450	331	2,631	2,922	2,867	2,398	539	143	54	34	1	1	1	2	
	PBCPC+	880	729	669	271	187	186	209	193	625	894	1,093	1,023	338	101	45	31	2	1	1	1	
	Biclique	-	-	-	-	-	2,375	1,985	11,003	-	-	-	-	-	29,402	3,606	355	-	25	61	93	
	BicliqueM	74	65	53	36	15	13	16	16	46	100	151	66	143	59	38	28	1	1	1	1	
	BicliqueP	79	45	36	28	16	12	12	11	35	38	46	36	188	64	40	29	2	1	1	1	

(we will explore the effectiveness of these partial-BCPCs in Exp-3).

PBCPC+ performs slightly better than MBAG and PBCPC, especially when α and β are small. For instance, on dataset Youtube, when $\alpha = \beta = 2$, PBCPC+ outperforms MBAG and PBCPC by more than one order of magnitude. Small α, β allow more maximal bicliques to satisfy the size threshold and become adjacent, making the MBAG larger. Thus, the MBAG compression brought by cross-subtree partial-BCPCs in PBCPC+ becomes more effective.

Biclique achieves relatively good performance only when both α and β are small—for example, on datasets Youtube, Bookcrossing, and DBLP with $\alpha = \beta = 2$. When α or β increases, the performance of Biclique drops sharply. This is because the number of (α, β) -bicliques grows exponentially as α and β increase.

BicliqueM and BicliqueP are the two optimal algorithms in almost all cases, and they can outperform MBAG by nearly three orders of magnitude. For example, on dataset Youtube, when $\alpha = \beta = 2$, BicliqueP only takes 34 seconds,

while MBAG requires 13,376 seconds. Furthermore, BicliqueP performs better than BicliqueM in most cases. This is because BicliqueP uses cross-subtree partial-BCPC for pruning, which has stronger pruning capability compared to maximal biclique (see Exp-4).

Exp-2: Overhead of computing partial-BCPC. This experiment evaluates the overhead of computing partial-BCPCs (with maximal biclique enumeration included, which is a necessary step for computing partial-BCPC). Specifically, we compare the time spent on computing partial-BCPCs in the algorithms PBCPC and PBCPC+, and also compare it against the time spent on maximal biclique enumeration only. The experimental results are shown in Figure 5.

As shown in Figure 5, although computing partial-BCPCs introduces additional overhead on top of maximal biclique enumeration, the added cost is manageable. On these datasets, t_{pbcpc} and $t_{\text{pbcpc+}}$ are slightly higher than the time spent on maximal biclique enumeration alone (t_{mbic}).

Exp-3: Effectiveness of partial-BCPC in partial-BCPC based solution. Table IV compares within-subtree partial-




Fig. 5. Time spent on maximal biclique enumeration (t_{mbic}), partial-BCPC computation (including maximal biclique enumeration as a necessary step) in PBCPC (t_{pbcpc}) and PBCPC+ (t_{pbcpc+})

TABLE IV

REMAINING RATIOS AFTER PARTIAL-BCPC COMPRESSION. n_{MB} DENOTES THE NUMBER OF MAXIMAL BICLIQUES (I.E., THE ORIGINAL MBAG VERTICES), n_{PBCPC} THE NUMBER OF PARTIAL-BCPCs IN PBCPC, AND n_{PBCPC+} THE NUMBER OF CROSS-SUBTREE PARTIAL-BCPCs IN PBCPC+

Dataset	α	β	n_{MB}	n_{PBCPC}	n_{PBCPC+}	$\frac{n_{PBCPC+}}{n_{MB}}$	$\frac{n_{PBCPC+}}{n_{PBCPC}}$
Youtube	2	2	1,769,331	292,951	48,346	2.7%	16.5%
	2	8	546,621	340,188	74,304	13.6%	21.8%
	8	2	768,396	373,311	107,580	14.0%	28.8%
	8	8	155,186	127,109	31,954	20.6%	25.1%
Bibsonomy	2	2	2,132,826	899,274	277,446	13.0%	30.9%
	2	8	799,631	353,784	186,738	23.4%	52.8%
	8	2	620,582	251,267	97,894	15.8%	39.0%
	8	8	371,569	211,478	97,237	26.2%	46.0%
Twitter	2	2	5,102,542	2,735,473	577,820	11.3%	21.1%
	2	8	460,594	323,463	135,860	29.5%	42.0%
	8	2	743,304	448,154	269,458	36.3%	60.1%
	8	8	108,721	79,833	35,196	32.4%	44.1%

BCPCs in PBCPC and cross-subtree partial-BCPCs in PBCPC+ as MBAG compression mechanisms on Youtube, Bibsonomy, and Twitter. Since each partial-BCPC is treated as one compressed MBAG vertex, we measure compression by the remaining vertex ratios n_{PBCPC+}/n_{MB} , and compare the two methods by n_{PBCPC+}/n_{PBCPC} . Across all settings, n_{PBCPC+}/n_{MB} is small (e.g., 2.7% on Youtube with $\alpha = \beta = 2$), and n_{PBCPC+}/n_{PBCPC} is well below 100% (e.g., 16.5% in the same setting and below 50% in most cases). Thus, cross-subtree merging provides substantial additional compression over within-subtree partial-BCPCs, explaining why PBCPC+ outperforms PBCPC in Table III.

Exp-4: Effectiveness of partial-BCPC in (α, β) -biclique enumeration based solution. Table V reports enumerated-node counts for Biclique, BicliqueM, and BicliqueP on Youtube, Bibsonomy, and Twitter. INF means Biclique does not finish within 24 hours and the ratio is “-”. Pruning is crucial: Biclique visits orders of magnitude more nodes and

TABLE V
NUMBER OF ENUMERATED NODES IN THE (α, β) -BICLIQUE ENUMERATION TREES. n_B , n_{BM} , AND n_{BP} DENOTE THE NUMBERS OF NODES ENUMERATED BY Biclique, BicliqueM, AND BicliqueP, RESPECTIVELY. INF INDICATES THAT THE ALGORITHM DID NOT FINISH WITHIN A 24-HOUR TIME BUDGET; THE CORRESPONDING RATIO ENTRIES ARE REPORTED AS “-”.

Dataset	α	β	n_B	n_{BM}	n_{BP}	$\frac{n_{BP}}{n_B}$	$\frac{n_{BP}}{n_{BM}}$
Youtube	2	2	1.4×10^7	788,775	114,679	0.817%	14.54%
	2	8	INF	1,887,512	433,791	-	22.98%
	8	2	INF	3,218,481	662,068	-	20.57%
	8	8	7.3×10^8	1,603,244	701,231	0.0964%	43.74%
Bibsonomy	2	2	6.4×10^8	2,063,463	597,505	0.0931%	28.96%
	2	8	INF	3,163,932	1,120,936	-	35.43%
	8	2	INF	2,143,285	527,585	-	24.62%
	8	8	INF	2,997,037	1,407,946	-	46.98%
Twitter	2	2	2.1×10^8	5,399,011	1,272,231	0.594%	23.56%
	2	8	INF	2,071,796	712,012	-	34.37%
	8	2	INF	3,652,114	1,665,074	-	45.59%
	8	8	INF	1,021,110	865,105	-	84.72%

TABLE VI
SCALABILITY TEST OF BicliqueP (s)

	Imdb	Actor2	Amazon	DBLP
20%	3.4	1.2	8.3	0.8
40%	7.8	5.7	54.5	1.4
60%	13.5	14.7	85.2	2.8
80%	17.5	25.8	107.3	4.5
100%	18.2	49.0	111.2	5.8

TABLE VII
MEMORY USAGE (MB) OF VARIOUS ALGORITHMS

	MBAG	PBCPC	PBCPC+	Biclique	BicliqueM	BicliqueP
Imdb	233	364	361	260	260	343
Actor2	409	725	720	449	449	639
Amazon	711	1,228	1,207	797	772	1,130
DBLP	935	957	970	1,200	935	940

often times out. For finished cases shown in the table, n_{BP}/n_B is always below 1% and reaches 0.0931% on Bibsonomy with $\alpha = \beta = 2$. Compared with BicliqueM, BicliqueP further reduces the search space, with n_{BP}/n_{BM} below 50% in nearly all shown settings. This explains the superior running time of BicliqueP in Table III.

Exp-5: Scalability. This experiment explores the scalability of our best algorithm, BicliqueP. Table VI presents the performance of BicliqueP on 4 large datasets. Specifically, we sampled 20%, 40%, 60%, and 80% of the edges from each dataset respectively. It can be observed that as the dataset size increases, the growth of running time of BicliqueP is manageable, demonstrating that BicliqueP has excellent scalability.

Exp-6: Memory. Table VII presents the memory usage of all algorithms on the largest 4 datasets. It can be observed that across all datasets, the 3 algorithms—MBAG, Biclique, and BicliqueM—consistently consume the least memory. This is because these algorithms do not require storing the MBE tree. However, although our best algorithm, BicliqueP, also stores the MBE tree, its memory consumption remains manageable.

Exp-7: Case study. We conduct a case study on the DBLP dataset. This dataset includes papers from major journals (TODS, TOIS, TKDE, VLDBJ) and conferences (SIGMOD, SIGKDD, ICDE, SIGIR, VLDB) in the database field, where vertices represent papers and authors respectively, and an edge (u, v) indicates that v is one of the authors of u . We compare (α, β) -BCPC with maximum biclique, k -bitruss [23], and (α, β) -core [24].

Given a bipartite graph $G = (U, V, E)$, (α, β) -core of G is a maximal subgraph of G , $G_1 = (U_1, V_1, E_1)$, that $\forall u \in$




Fig. 6. Case study on DBLP

$U_1, |Nei_{G'}(u)| \geq \alpha, \forall v \in V_1, |Nei_{G'}(v)| \geq \beta$. k -bitruss is a maximal subgraph of G , $G_2 = (U_2, V_2, E_2)$, that each edge in E_2 is contained in at least k butterflies (butterfly is a subgraph with four vertices: $u, w \in U_2, v, x \in V_2, (u, v) \in E_2, (u, x) \in E_2, (w, v) \in E_2, (w, x) \in E_2$).

In this case study, $\alpha = 4, \beta = 4, k = 12$. By Definition 4, the smallest maximal bicliques in $(4, 4)$ -BCPC are either $(5, 4)$ -bicliques or $(4, 5)$ -bicliques, and the edges within these maximal bicliques exist in at least $3 \times 4 = 12$ butterflies, thus, we set $k = 12$ for k -bitruss.

Figure 6 presents the results containing Jiawei Han across all models. Among them, the result from maximum biclique is the smallest. Compared with $(4, 4)$ -BCPC, it lacks one author, Wei Fan, who is one of Jiawei Han's collaborators on Google Scholar. Although 12-bitruss and $(4, 4)$ -core include the results of $(4, 4)$ -BCPC, they also contain a large number of other vertices, making the usability of these two models far inferior to that of $(4, 4)$ -BCPC.

VII. RELATED WORK

BCPC and its applications. [7] proposed the BCPC model and a detection method based on traversing the MBAG; [15] implemented this as an online toolkit. Applications include user-cluster mining in social networks [11], editor-article relationship analysis in Wikipedia [8], abnormal-behavior detection in enterprise networks [9], and student-resource pattern analysis in online learning [10], [12]. [13], [14] further studied personalized BCPC search.

Maximal biclique and (p, q) -biclique enumeration. Maximal biclique enumeration underpins BCPC detection. The powerset-based framework of [25] was improved by pivot pruning [19], optimized vertex ordering [20], and a unified hereditary-subgraph framework [16]. For (p, q) -biclique enumeration, [26] uses it as a subroutine for broader analysis, while [17] focuses on optimizing the task itself; [18] addresses exact and approximate counting.

Other cohesive subgraph models. Related models include (α, β) -core [27], [24], [28], k -bitruss [29], [23], [30], and maximum biclique [31], [32]. Modularity-based [33], [34] and label-propagation [35] methods are also relevant.

VIII. CONCLUSION

In this work, we study biclique percolation community (BCPC) detection in bipartite networks. To address the massive maximal biclique adjacency graph (MBAG) used by existing methods, we introduce cross-subtree partial-BCPCs to group maximal bicliques that provably belong to the same BCPC even when they are enumerated in different subtrees, thereby reducing redundant MBAG traversal. We further propose an (α, β) -biclique enumeration method that connects

maximal bicliques sharing enumerated (α, β) -bicliques, with cross-subtree partial-BCPCs pruning the enumeration space. Experiments show that our methods outperform existing approaches by nearly three orders of magnitude, enabling more scalable BCPC detection on large bipartite networks.

REFERENCES

- [1] S. Fortunato, "Community detection in graphs," *Physics reports*, vol. 486, no. 3-5, pp. 75–174, 2010.
- [2] K. Steinhauser and N. V. Chawla, "Community detection in a large real-world social network," in *Social computing, behavioral modeling, and prediction*. Springer, 2008, pp. 168–175.
- [3] S. Rahiminejad, M. R. Maurya, and S. Subramaniam, "Topological and functional comparison of community detection algorithms in biological networks," *BMC bioinformatics*, vol. 20, no. 1, p. 212, 2019.
- [4] C. Choudhary, I. Singh, and M. Kumar, "Community detection algorithms for recommendation systems: techniques and metrics," *Computing*, vol. 105, no. 2, pp. 417–453, 2023.
- [5] C. Maier and D. Simovici, "Bipartite graphs and recommendation systems," *Journal of Advances in Information Technology-in print*, 2022.
- [6] C. Carusi and G. Bianchi, "A look at interdisciplinarity using bipartite scholar/journal networks," *Scientometrics*, vol. 122, no. 2, pp. 867–894, 2020.
- [7] S. Lehmann, M. Schwartz, and L. K. Hansen, "Biclique communities," *Physical Review E—Statistical, Nonlinear, and Soft Matter Physics*, vol. 78, no. 1, p. 016108, 2008.
- [8] R. Jesus, M. Schwartz, and S. Lehmann, "Bipartite networks of wikipedia's articles and authors: a meso-level approach," in *Proceedings of the 5th international symposium on Wikis and open collaboration*, 2009, pp. 1–10.
- [9] Q. Liao, A. Striegel, and N. Chawla, "Visualizing graph dynamics and similarity for enterprise network security and management," in *Proceedings of the seventh international symposium on visualization for cyber security*, 2010, pp. 34–45.
- [10] T. Hecking, S. Ziebarth, and H. U. Hoppe, "Analysis of dynamic resource access patterns in a blended learning course," in *Proceedings of the Fourth International Conference on Learning Analytics and Knowledge*, 2014, pp. 173–182.
- [11] T. Hecking, I.-A. Chounta, and H. U. Hoppe, "Analysis of user roles and the emergence of themes in discussion forums," in *2015 Second European Network Intelligence Conference*. IEEE, 2015, pp. 114–121.
- [12] S. Ziebarth, G. Neubaum, E. Kyewski, N. Kramer, H. U. Hoppe, T. Hecking, and S. Eimler, "Resource usage in online courses: Analyzing learner's active and passive participation patterns." International Society of the Learning Sciences, Inc.[ISLS], 2015.
- [13] Z. Chen, Y. Zhao, L. Yuan, X. Lin, and K. Wang, "Index-based biclique percolation communities search on bipartite graphs," in *2023 IEEE 39th International Conference on Data Engineering (ICDE)*. IEEE, 2023, pp. 2699–2712.
- [14] Z. Chen, H. Huang, L. Yuan, J. Xu, B. Zheng, and X. Lin, "Biclique percolation communities computation on temporal bipartite graphs," *IEEE Transactions on Knowledge and Data Engineering*, 2026.
- [15] B. Wang, J. Chen, and S. Zhang, "Bmtk: a toolkit for determining modules in biological bipartite networks," *Quantitative Biology*, vol. 6, pp. 186–192, 2018.
- [16] Q. Dai, R.-H. Li, X. Ye, M. Liao, W. Zhang, and G. Wang, "Hereditary cohesive subgraphs enumeration on bipartite graphs: The power of pivot-based approaches," *Proceedings of the ACM on Management of Data*, vol. 1, no. 2, pp. 1–26, 2023.
- [17] J. Yang, Y. Peng, D. Ouyang, W. Zhang, X. Lin, and X. Zhao, " (p, q) -biclique counting and enumeration for large sparse bipartite graphs," *The VLDB Journal*, vol. 32, no. 5, pp. 1137–1161, 2023.
- [18] X. Ye, R.-H. Li, Q. Dai, H. Qin, and G. Wang, "Efficient biclique counting in large bipartite graphs," *Proceedings of the ACM on Management of Data*, vol. 1, no. 1, pp. 1–26, 2023.
- [19] A. Abidi, R. Zhou, L. Chen, and C. Liu, "Pivot-based maximal biclique enumeration," in *IJCAI*, 2020, pp. 3558–3564.
- [20] L. Chen, C. Liu, R. Zhou, J. Xu, and J. Li, "Efficient maximal biclique enumeration for large sparse bipartite graphs," *Proceedings of the VLDB Endowment*, vol. 15, no. 8, pp. 1559–1571, 2022.
- [21] T. H. Cormen, C. E. Leiserson, R. L. Rivest, and C. Stein, *Introduction to algorithms*. MIT press, 2022.
- [22] E. Prisner, "Bicliques in graphs i: Bounds on their number," *Combinatorica*, vol. 20, no. 1, pp. 109–117, 2000.

- [23] K. Wang, X. Lin, L. Qin, W. Zhang, and Y. Zhang, "Efficient bitruss decomposition for large-scale bipartite graphs," in *2020 IEEE 36th International Conference on Data Engineering (ICDE)*. IEEE, 2020, pp. 661–672.
- [24] B. Liu, L. Yuan, X. Lin, L. Qin, W. Zhang, and J. Zhou, "Efficient (α, β) -core computation in bipartite graphs," *The VLDB Journal*, vol. 29, no. 5, pp. 1075–1099, 2020.
- [25] Y. Zhang, C. A. Phillips, G. L. Rogers, E. J. Baker, E. J. Chesler, and M. A. Langston, "On finding bicliques in bipartite graphs: a novel algorithm and its application to the integration of diverse biological data types," *BMC bioinformatics*, vol. 15, no. 1, p. 110, 2014.
- [26] M. Mitzenmacher, J. Pachocki, R. Peng, C. Tsourakakis, and S. C. Xu, "Scalable large near-clique detection in large-scale networks via sampling," in *Proceedings of the 21th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2015, pp. 815–824.
- [27] B. Liu, L. Yuan, X. Lin, L. Qin, W. Zhang, and J. Zhou, "Efficient (α, β) -core computation: An index-based approach," in *The World Wide Web Conference*, 2019, pp. 1130–1141.
- [28] Q. Liu, X. Liao, X. Huang, J. Xu, and Y. Gao, "Distributed (α, β) -core decomposition over bipartite graphs," in *2023 IEEE 39th International Conference on Data Engineering (ICDE)*. IEEE, 2023, pp. 909–921.
- [29] Z. Zou, "Bitruss decomposition of bipartite graphs," in *International conference on database systems for advanced applications*. Springer, 2016, pp. 218–233.
- [30] F. Lin, B. Ruan, J. Gan, and L. Li, "Efficient bitruss decomposition without butterfly enumeration," in *Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining V. 2*, 2025, pp. 1718–1728.
- [31] B. Lyu, L. Qin, X. Lin, Y. Zhang, Z. Qian, and J. Zhou, "Maximum biclique search at billion scale," *Proceedings of the VLDB Endowment*, 2020.
- [32] K. Wang, W. Zhang, X. Lin, L. Qin, and A. Zhou, "Efficient personalized maximum biclique search," in *2022 IEEE 38th International Conference on Data Engineering (ICDE)*. IEEE, 2022, pp. 498–511.
- [33] M. J. Barber, "Modularity and community detection in bipartite networks," *Physical Review E—Statistical, Nonlinear, and Soft Matter Physics*, vol. 76, no. 6, p. 066102, 2007.
- [34] R. Arthur, "Modularity and projection of bipartite networks," *Physica A: Statistical Mechanics and its Applications*, vol. 549, p. 124341, 2020.
- [35] X. Liu and T. Murata, "An efficient algorithm for optimizing bipartite modularity in bipartite networks," *Journal of Advanced Computational Intelligence and Intelligent Informatics*, vol. 14, no. 4, pp. 408–415, 2010.